

# *Introduction to Database Design*



# This course

After this class, you should be able to...

- Use a relational database (via SQL and code)
- Design a secure, normalized, efficient database
- Understand (some of) how a DBMS works
- Understand database transactions, security and recovery.

Expectations from you

- a) Work hard (really hard, it'll be worth it!)
- b) Use resources (read book/materials before class, attend class, etc.)
- c) Always complete the labs/assignments

# Course Overview



Week No.	Database Design
1	Introduction to Database Design
2	Conceptual Database Design
3	Logical Design
4	Normalization
5	SQL& Relational algebra part 1
6	SQL& Relational algebra part 2
7	MID TERM
8	SQL and Web Programming
9	SQL Query Optimizations <b>Set Exercises (70%)</b> <b>AE1 Submission 14/11/2025</b>
10	Database Security and Administration
11	Database Transactions and Concurrency Part I
12	Autumn Break
13	Database Transactions and Concurrency Part I <b>Written Assignment(Group &amp; Individual)</b> <b>AE2 Submission 8/12/2025</b>

# labs

- Labs/Assignments & Feedbacks
- Will help with AE1 & AE 2
- Will be used for check-in on the group assignments

# Outline

1. What is a Database? A DBMS?
2. Why use a DBMS?
3. Databases in Context
4. Design and Implementation Process
5. Relational Data Model (*Intro*)
6. *DBMS in the Real World - Industry Research & Presentation (lab)*
7. *Software installation(lab)*

# What is Database?

A collection of related data, most often...

- reflects some aspect of the real world
- logically coherent with inherent meaning
- designed, built, and populated with data for a specific purpose
  - intended group of users
  - some preconceived applications with which these users are interested
  - application requirements in terms of performance, security, redundancy, concurrency, etc.

# Database Management System (DBMS)

A collection of programs that enables users to create and maintain a database

- Supports specifying the data types, structures, and constraints of the data
- Stores the data on some medium under control of the DBMS
- Supports querying and updating the database
- Protects data against malfunction and unauthorized access

# Why use a DBMS?

Common tradeoff in CS:

A. Code from scratch

- Pros: you know your problem best (so fast, customized)
- Cons: slow, labor intensive, need to add/change features?

B. Find a library/tool that solves [part of] your problem

- Pros: fast via bootstrapping, better designed?
- Cons: understand the tool, may not be efficient, support?

DBMSs adopt some set of limiting assumptions in order to efficiently support a useful feature set over a wide class of possible databases



# Example: Student Records

- Given a school with MANY students (NEU: ~25k, UM: ~45k), each with some data (name, ID, DOB, classes)
- Write a program that can efficiently...
  - Retrieve a random student
  - Retrieve the first/last student, according to...
    - Last name
    - DOB
  - Retrieve a student by...
    - ID
    - Name (with \*'s)
  - Retrieve a class roster (all students in class X)
  - Handles adding/removing/editing students/classes
  - Handles multiple simultaneous reads/writes
  - Provides differing access rights
  - Handles OS faults/power outages

...

# Many Kinds of DBMSs (1)

- Graph databases
  - Create nodes, edges, labels
  - Query about relationships and paths
    - Find your friends
    - Find someone that can help you learn databases



- Spatial databases
  - Data pertaining to space occupied by objects
  - Objects in 2D/3D
  - Query locations, relations
    - Collision detection



# Many Kinds of DBMSs (2)

- Key-Value stores
  - Associative array
  - Scalable, fault-tolerant
  - Query



Key	Value
K1	AAA,BBB,CCC
K2	AAA,BBB
K3	AAA,DDD
K4	AAA,2,01/01/2015
K5	3,ZZZ,5623

- Document stores
  - Create dynamic documents
  - Query about contents
    - Find by author, title, content, etc. patterns





# Relational DBMS

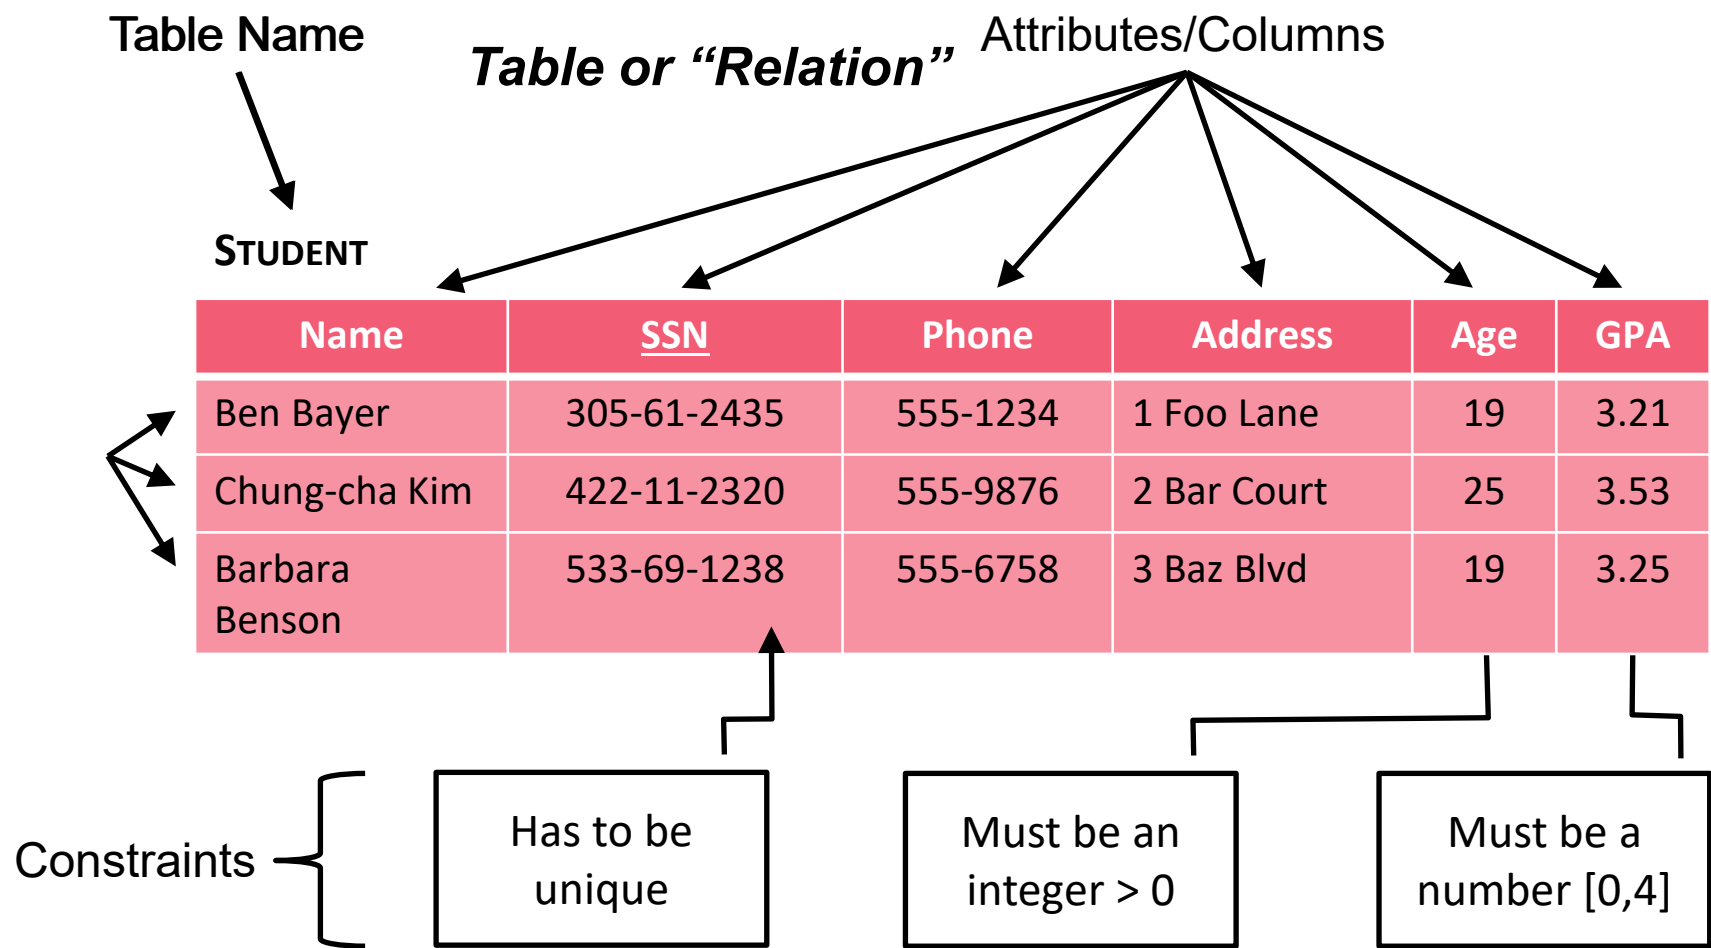
We focus on **relational** databases

Based on the relational *data model*


- Researched ~45 years, widely used
  - Free/paid implementations for personal use, embedded systems, small/large enterprise



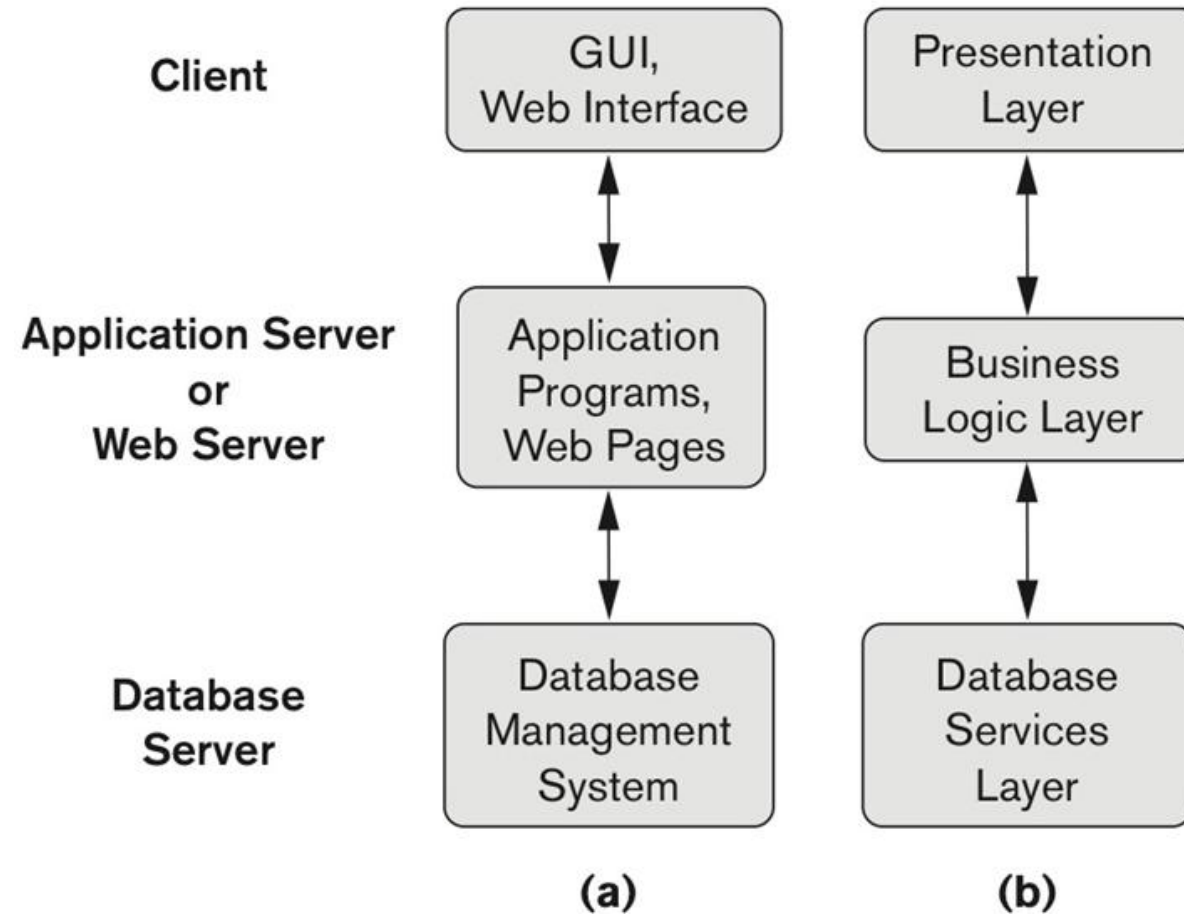
# Relational Databases (1)



# Relational Database (2)

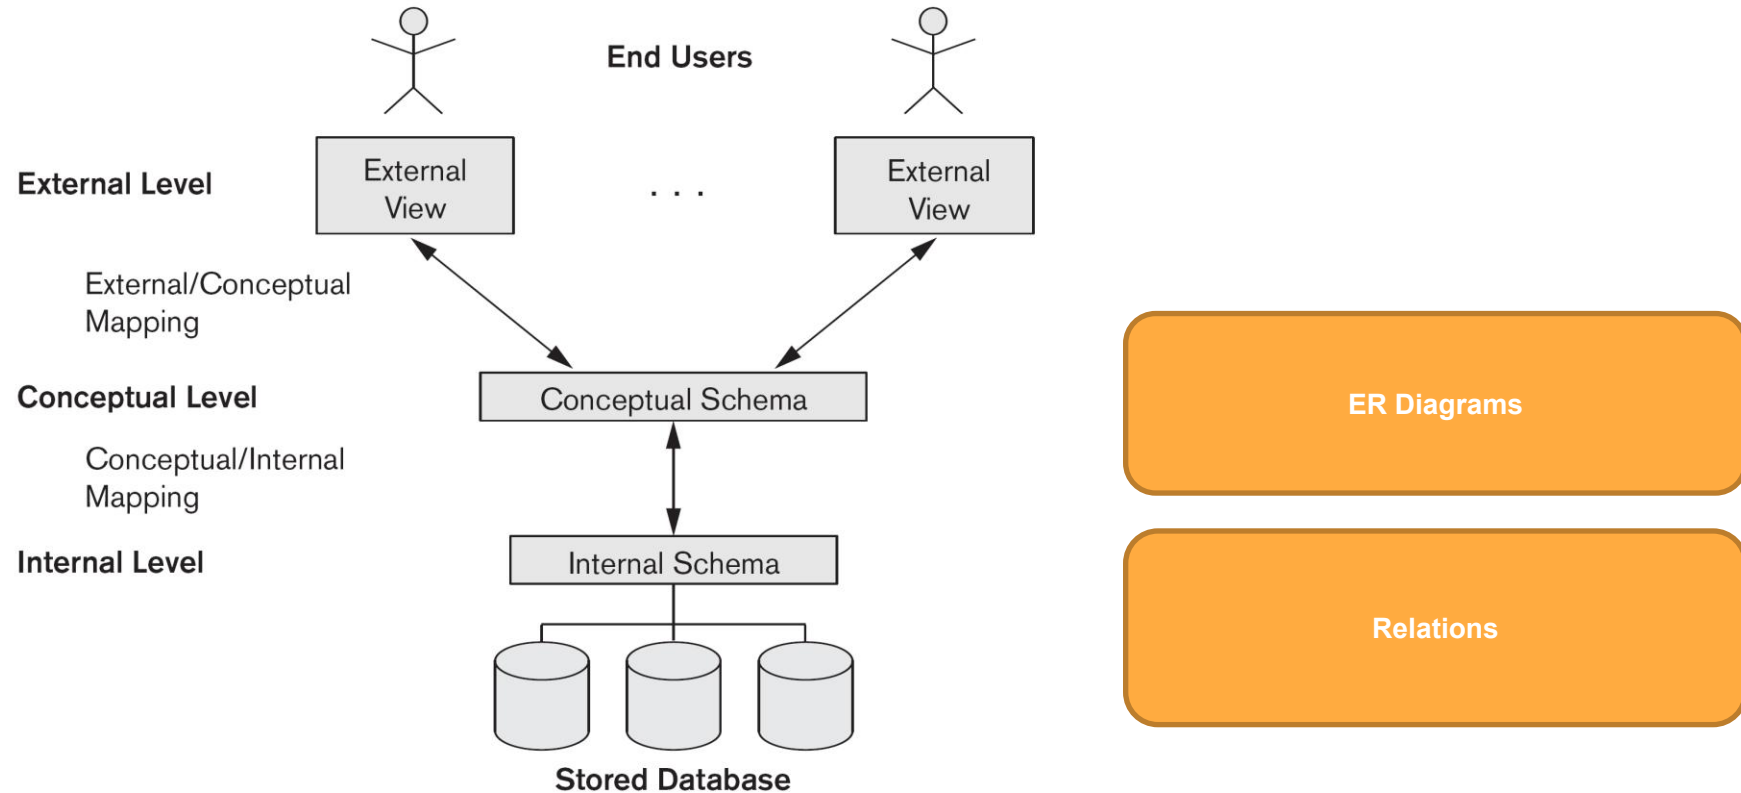


# Databases in Context



# Relational Database features (1)

- Data *independence* via **data models**
  - Conceptual representation independent of underlying storage or operation implementation





# Relational DBMS Features (2)

- Operation abstraction via...
  - Declarative languages
    - Structured Query Language (SQL)
      - Data... definition, manipulation, query
  - Programmatic APIs
    - Function libraries (focus), embedded languages, stored procedures, etc.

# Relational DBMS Features (3)

- Reliable concurrent transactions
  - **(A)**tomocity: “all or nothing”
  - **(C)**onsistency: valid -> valid’
  - **(I)**solation: parallel execution, serial result
  - **(D)**urability: once it is written, it is so
- High performance
  - Buffering, caching, locking (like a mini OS)
  - Query optimization, redundant data structures (e.g. indexes, materialized views)

# Relational DBMS Features (4)

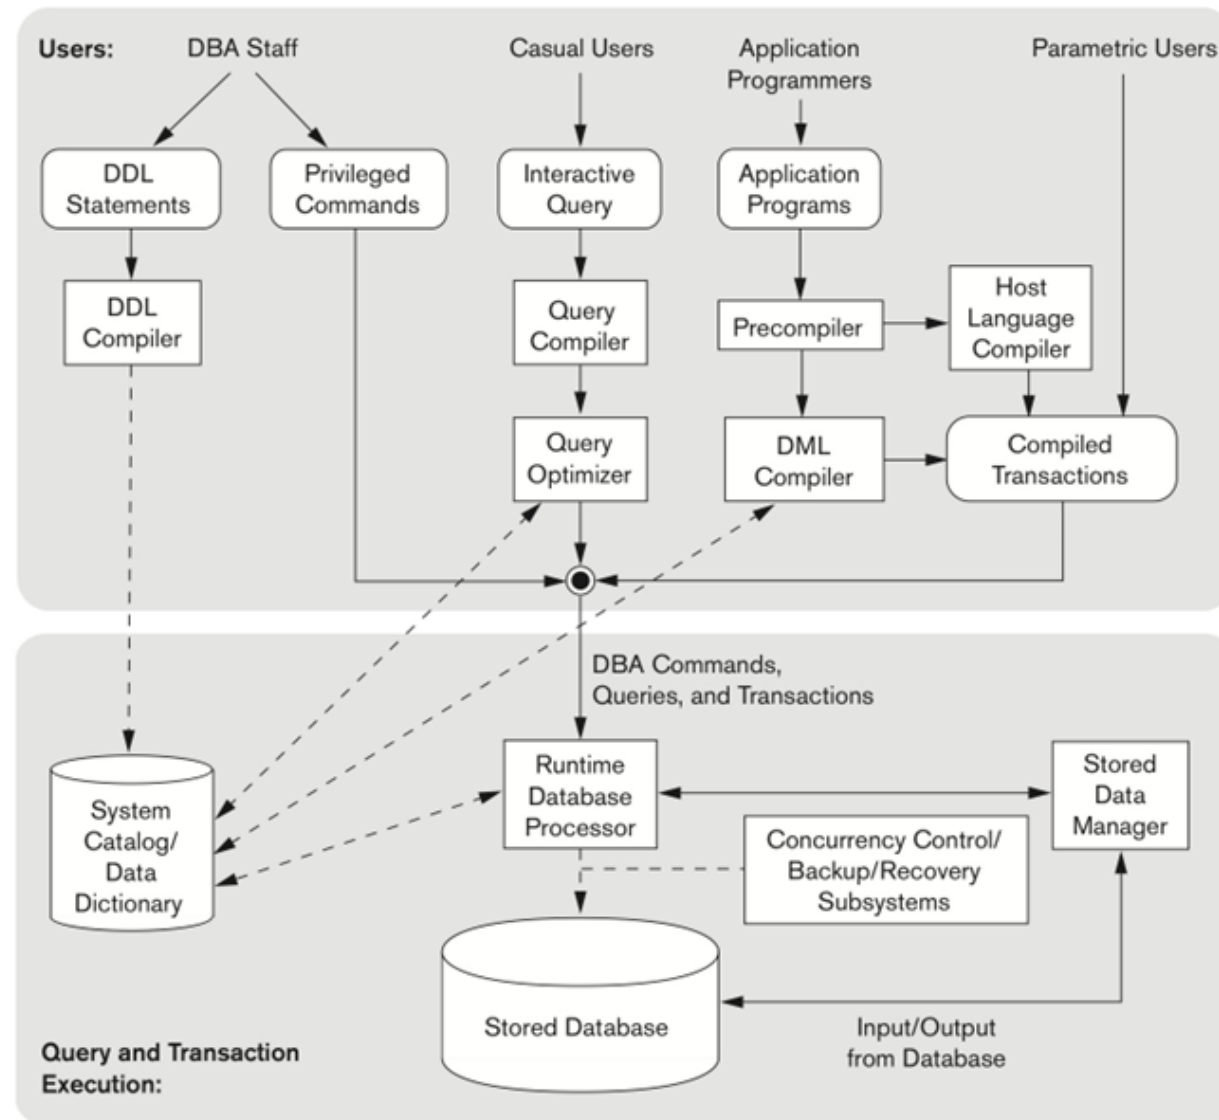
- Authentication and authorization
  - Discussed in context of other security concerns/techniques
- Backup and recovery
  - Logging, replication, migration

# Databases in Context

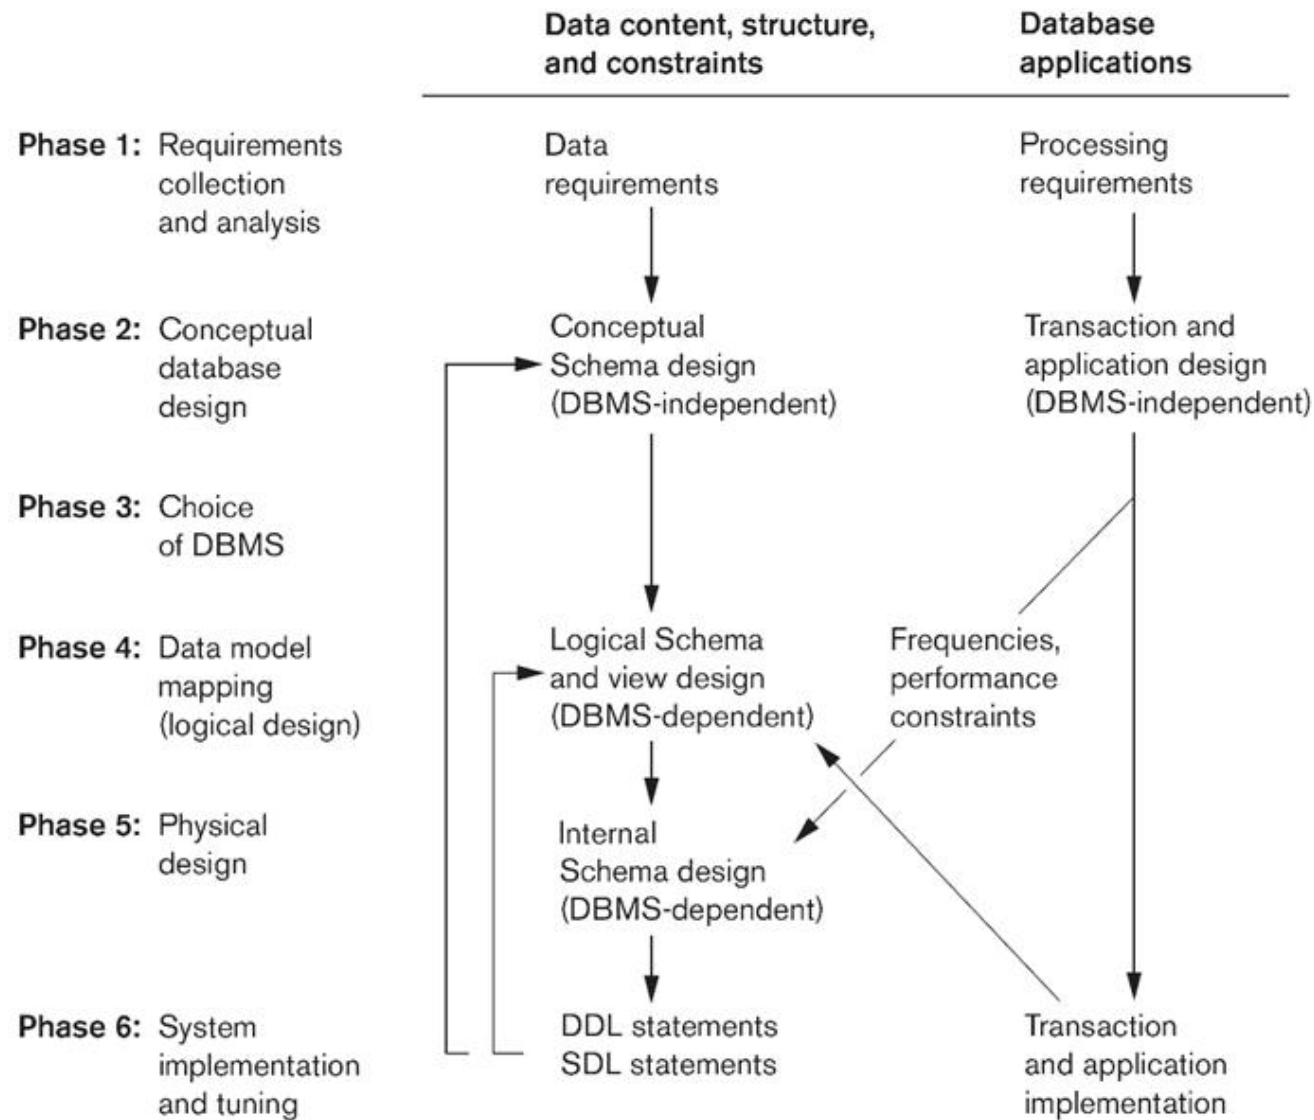
## *People*

1. **Database designers**
2. **System analysts & application programmers**
3. Database administrators
4. End users
5. Back-end
  - a. DBMS designer/implementer
  - b. Tool developers
  - c. SysAdmins

# Relational DBMS



# Database Design and Implementation Process



# Requirements Collection & Analysis

- Data/Constraints

*“The company is organized into departments. Each department has a unique name, number, and a particular employee who manages the department. We keep track...”*

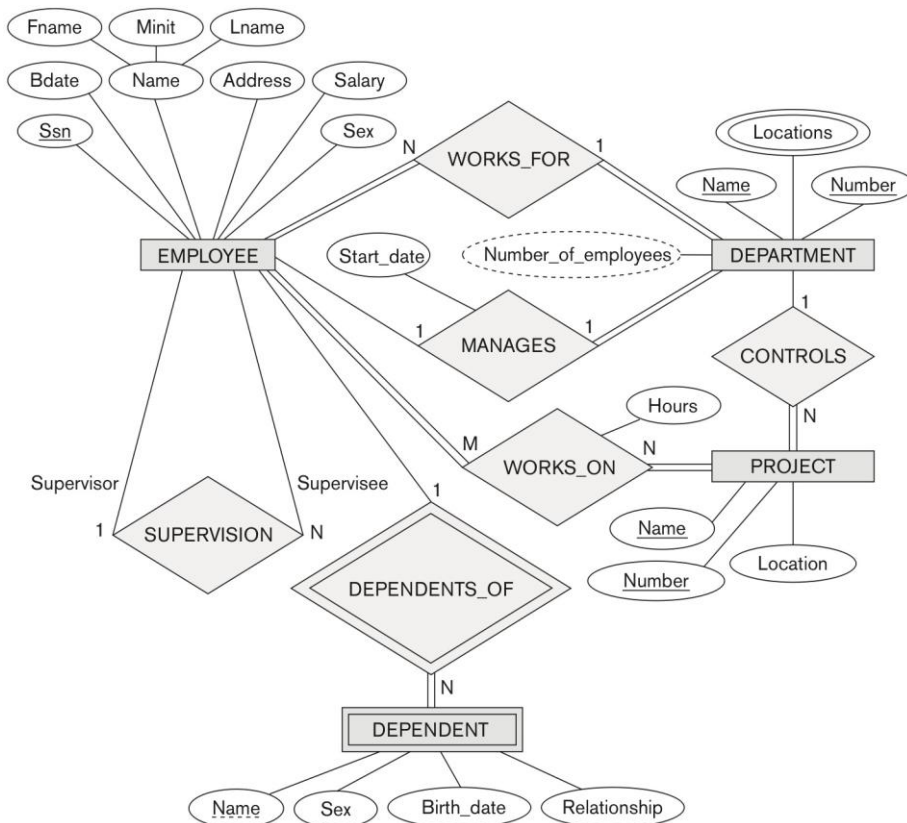
- Functional Needs

- Operations/queries/reports
  - Frequency
- Performance, security, etc.



# Conceptual Design

## Data



## Application

- Software
  - UML
  - Form design
- Database
  - Transaction design
  - Report design



# Logical Design

## Data

### EMPLOYEE

Fname	Minit	Lname	<u>Ssn</u>	Bdate	Address	Sex	Salary	Super_ssn	Dno
-------	-------	-------	------------	-------	---------	-----	--------	-----------	-----

### DEPARTMENT

Dname	<u>Dnumber</u>	Mgr_ssn	Mgr_start_date
-------	----------------	---------	----------------

### DEPT\_LOCATIONS

<u>Dnumber</u>	<u>Dlocation</u>
----------------	------------------

### PROJECT

Pname	<u>Pnumber</u>	Plocation	Dnum
-------	----------------	-----------	------

### WORKS\_ON

<u>Essn</u>	<u>Pno</u>	Hours
-------------	------------	-------

### DEPENDENT

<u>Essn</u>	<u>Dependent_name</u>	Sex	Bdate	Relationship
-------------	-----------------------	-----	-------	--------------

## Application

- Supporting code (that does not depend upon database)

# Normalization

## Data

- Schema Refinement

## Application

- Supporting code (that does not depend upon database)

# Physical Design

## Data

- Index, materialized view selection and analysis

## Application

- Implementing operations as queries
- Implementing constraints as keys, triggers, views
- Implementing multi-user security as grants

# Implementing and Tuning

## Data

- DDL statements
- De-normalization, updating indexes/materialized views

## Application

- Query integration
- Profiling queries/operations
- Security, concurrency, performance, etc. analysis

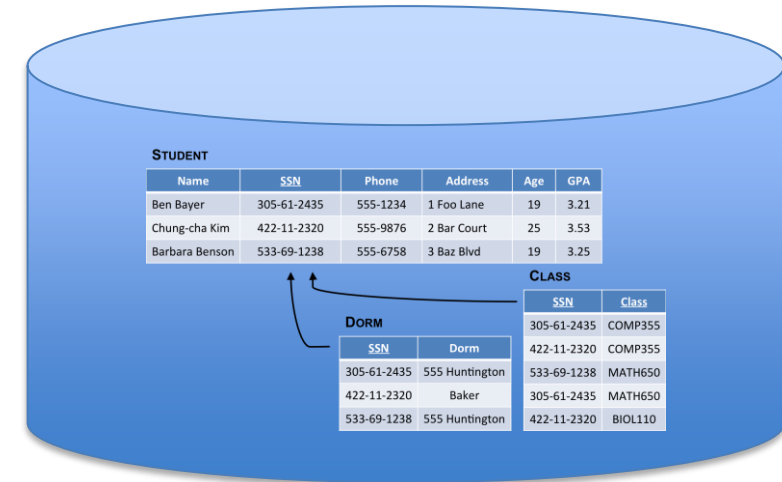
# Relational Data Model

A database consists of...

- i. a set of ***relations*** (tables)
- ii. a set of ***integrity constraints***

A database is in a **valid state** if it satisfies all integrity constraints (else **invalid state**)

Pop Quiz:  
What is a set?



# A Relation

A relation consists of...

- i. its ***schema***, describing structure
- ii. its ***state***, or current populated data

Schema	STUDENT					
	Name	<u>SSN</u>	Phone	Address	Age	GPA
State	Ben Bayer	305-61-2435	555-1234	1 Foo Lane	19	3.21
	Chung-cha Kim	422-11-2320	555-9876	2 Bar Court	25	3.53
	Barbara Benson	533-69-1238	555-6758	3 Baz Blvd	19	3.25

# Relational Schema

- Relation name  
STUDENT
- Ordered list of  $n$  **attributes** (columns; degree  $n$  or  $n$ -ary)  
Each with a corresponding **domain** (set of valid **atomic** values)  
dom(SSN) = “###-##-####”  
dom(GPA) = [0, 4]

Notation: NAME( $A_1, A_2, \dots A_n$ )

STUDENT(Name, SSN, Phone, Address, Age, GPA)

What is the degree of  
STUDENT?

**STUDENT**

Name	<u>SSN</u>	Phone	Address	Age	GPA
------	------------	-------	---------	-----	-----

# Relation State

A set of *n*-**tuples** (rows)

Each has a value in the domain of every corresponding attribute (or **NULL**)

Notation:  $r(\text{NAME})$

Mathematically, a subset of the Cartesian product of the attribute domains; related to the closed-world assumption

$$r(STUDENT) \subseteq (dom(Name) \times dom(SSN) \times \dots dom(GPA))$$

Ben Bayer	305-61-2435	555-1234	1 Foo Lane	19	3.21
Chung-cha Kim	422-11-2320	555-9876	2 Bar Court	25	3.53
Barbara Benson	533-69-1238	555-6758	3 Baz Blvd	19	3.25



## Exercise

Diagrammatically produce a relation HAT according to the following schema; the relation state should have at least three tuples

HAT(Team, Size, Color)

$\text{dom}(\text{Team}) = \{ \text{RedSox}, \text{Bruins}, \text{Celtics}, \text{Patriots}, \text{Revolution} \}$

$\text{dom}(\text{Size}) = \{ S, M, L, XL \}$

$\text{dom}(\text{Color}) = \{ \text{Black}, \text{Blue}, \text{White}, \text{Red}, \text{Green}, \text{Yellow} \}$

How many tuples are possible in this relation?

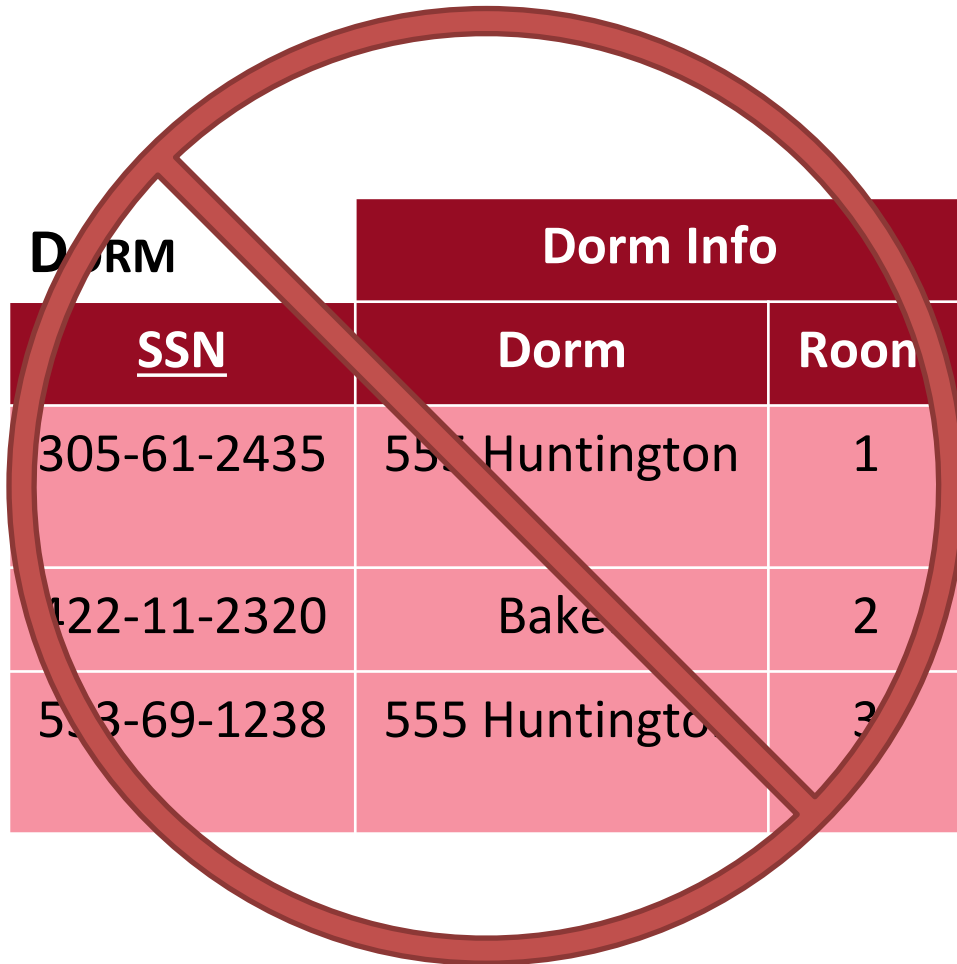
# NULL

- **NULL** is a special value that may be in the attribute domain
- Used to represent the values of attributes that may be unknown or may not apply to a tuple.
- Several possible meanings
  - e.g. unknown, not available, does not apply, undefined, ...
- Best to avoid
- Else deal with caution

# Value Structure in Tuples

- Each value should be **atomic** – no *composite* or *multi-valued* attributes
  - Composite: “one column, many parts”
  - Multi-valued: “one column, multiple values”
- Convention called 1NF (*first normal form*)
  - More on this later in the course

# Violation of 1NF: Composite



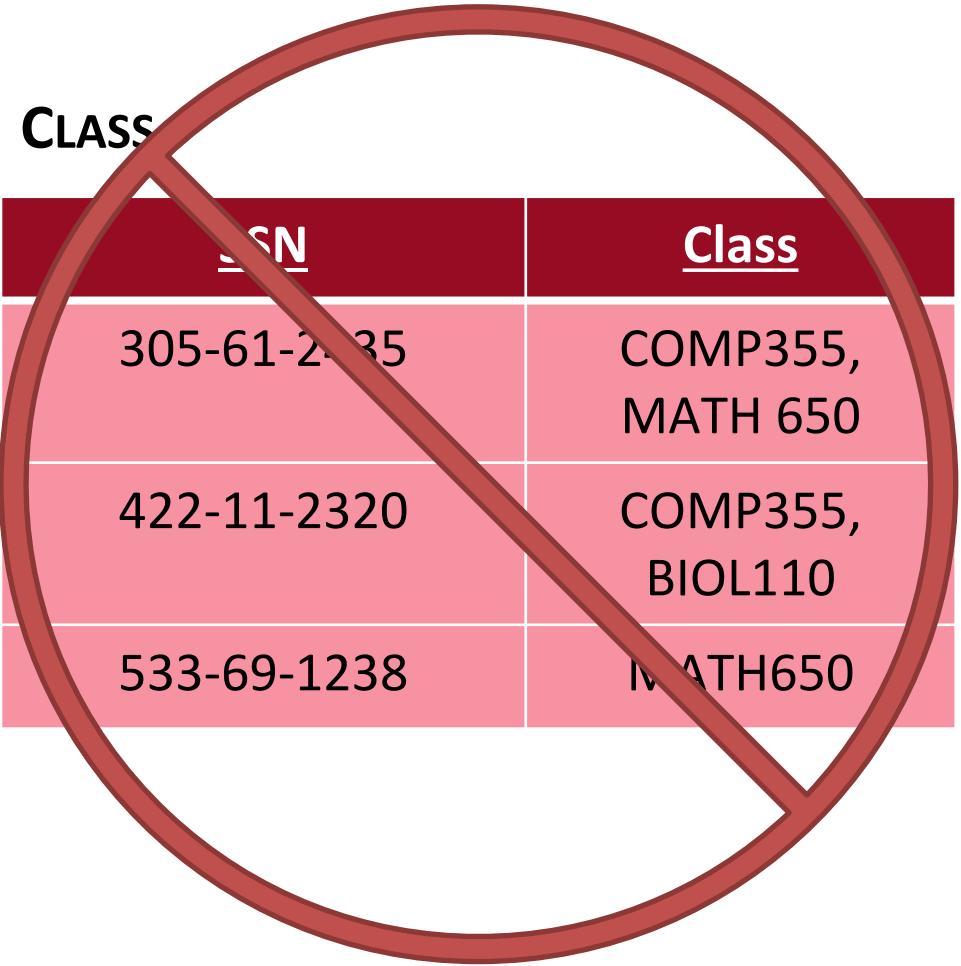
DORM	Dorm Info	
	<u>SSN</u>	Room
	305-61-2435	555 Huntington 1
	422-11-2320	Baker 2
	533-69-1238	555 Huntington 3

vs.

DORM		
<u>SSN</u>	Dorm	Room
305-61-2435	555 Huntington	1
422-11-2320	Baker	2
533-69-1238	555 Huntington	3

# Violation of 1NF: Multi-Valued

CLASS



<u>SSN</u>	<u>Class</u>
305-61-2435	COMP355, MATH 650
422-11-2320	COMP355, BIOL110
533-69-1238	MATH650

**vs.**

CLASS

<u>SSN</u>	<u>Class</u>
305-61-2435	COMP355
422-11-2320	COMP355
533-69-1238	MATH650
305-61-2435	MATH650
422-11-2320	BIOL110

# Summary

- A **database** is a collection of related data that reflects some aspect of the real world; is logically coherent with inherent meaning; and is designed, built, and populated with data for a specific purpose
- A **database management system** (DBMS) is a collection of programs that enables users to create and maintain a database
- There are many types – we will focus on **relational** databases (RDBMS) and a bit of NoSQL Databases
- The typical database design process is an iterative process of requirements collection/analysis, conceptual design, logical design, physical design, and system implementation/tuning

# Friday Lab

- The lab is available on canvas
  - Please try to go through before Friday.

# Acknowledgement

- The major part of the teaching agenda for this course is based on material developed by Nate Derbinsky (<https://derbinsky.info/>)